

KaryoFlow: Few-Step Rectified-Flow Chromosome Detection With Localization-Quality Calibration

Author Name, *Member, IEEE*, and Co-Author Name, *Senior Member, IEEE*

Abstract—Automated karyotyping must detect about 46 densely distributed chromosomes from 24 morphologically similar classes. Small extent, touching or overlap, weak boundaries, and near-saturated class confidence make precise localization and reliable ranking the main obstacles to high-accuracy detection. We present *KaryoFlow*, which addresses these task-specific difficulties through low-dimensional rectified box transport, identity-consistent refinement of adjacent proposals, and final localization-quality calibrated ranking. On two public chromosome datasets with markedly different object-scale distributions, *KaryoFlow* exceeds the strongest evaluated detector by relative mAP margins of 2.2% and 0.2%, respectively. Combining few-step inference with cascade-head distillation provides up to $2.73\times$ faster inference than the RF+Heun reference, while distillation reduces mAP by only 0.0003 relative to its six-head teacher. These results establish a localization-aware generative approach and provide a new direction for high-precision chromosome detection.

Index Terms—Rectified Flow, object detection, localization quality, diffusion models, medical image analysis, chromosome karyotyping

I. INTRODUCTION

CHROMOSOME karyotyping remains a labor-intensive step in cytogenetic analysis. Its detection stage is structurally unusual: a metaphase spread contains a nearly fixed set of about 46 chromosomes assigned to 24 classes, many instances occupy less than one percent of the image, and neighboring chromosomes frequently overlap in their bounding-box support. In our two held-out cohorts, the median instance count is 46; 99.1% and 100% of images contain at least one pair of ground-truth boxes with IoU above 0.1. The median relative box area also shifts from 0.4% to 0.8% between cohorts. Consequently, a few-pixel error can preserve AP50 while changing the strict-IoU ordering that determines mAP, and an iterative detector must refine many adjacent candidates without confusing their identities.

Conventional detectors offer complementary but incomplete solutions. Anchor-based systems such as Cascade R-CNN [1] and YOLOX [2] are fast, whereas transformer detectors such as DINO [3] obtain strong localization through multi-scale deformable attention [4]. DiffusionDet [5] instead treats boxes as a set generated from noise, an appealing formulation for variable spatial layouts, but its DDPM inheritance introduces a curved stochastic path and a many-step design that are poorly




Fig. 1. Speed-accuracy comparison on two chromosome cohorts. Accuracy follows the locked evaluation protocol; model-forward FPS is measured with the same 512×512 , batch-1 RTX A6000 protocol. Accuracy points in both panels are three-inference-seed means from fixed checkpoints. Both panels include three unconnected RF-DPM++ operating points at $K = 500$, 200, and 100; the Dataset 1 $K = 200/100$ points use inference-only pruning of the same checkpoint. Higher and farther right is better.

matched to interactive microscopy analysis. More importantly, simply shortening that path does not resolve two detection-specific mismatches: multistep integration assumes a persistent state identity, whereas proposal renewal replaces selected states; and final ranking by class confidence ignores whether a small chromosome is localized accurately enough to survive strict IoU thresholds.

Rectified Flow [6] (RF) offers a natural generation foundation by replacing the DDPM path with a deterministic linear transport between noise boxes and targets. Applying RF to

detection raises three linked questions: how to learn a well-conditioned box trajectory from limited data, how to integrate it without violating proposal identity, and how to rank fixed outputs by both semantic and geometric reliability. *KaryoFlow* answers them as one trajectory-decision design, shown in Figure 2. Its contributions are threefold.

First, we formulate RF directly in normalized four-dimensional box space and derive a data-prediction objective

Manuscript received [date]; revised [date]. (Corresponding author: Author Name.)

The authors are with [Affiliation], [City], [Country] (e-mail: author@example.com).

for the reverse path. Although data and velocity prediction are algebraically equivalent, their squared losses differ by t^{-2} ; direct box prediction therefore avoids amplifying errors near the terminal state emphasized by the shifted schedule. RF itself is adopted, and our claim is the detection formulation and analysis rather than invention of the underlying flow framework.

Second, we derive an RF data-prediction update from DPM-Solver++ [7] and establish the proposal-identity condition for multistep box integration. If a row is renewed between two evaluations, the history correction becomes a secant between different candidates and is not a derivative along either trajectory. Disabling renewal when the proposal pool is redundant restores this premise and saves 2.43% latency; its failure at $K = 100$ provides an explicit capacity boundary. DPM-Solver++ itself is used for efficiency, not claimed as an accuracy novelty.

Third, we introduce Localization-Quality Calibrated Ranking (LQCR) for the strict-IoU bottleneck of small, crowded chromosomes. From the final proposal feature, LQCR predicts conditional localization quality and ranks detections by $s = pq^2$. It is applied only when outputs are emitted, leaving candidate coordinates, class logits, proposal survival, and the entire RF trajectory unchanged before final post-processing. The resulting fixed-checkpoint gain is therefore attributable to decision-stage ranking.

We evaluate KaryoFlow on two public chromosome datasets containing 2,200 and 5,000 images, with multi-seed controls, same-checkpoint solver tests, test-set evaluation, strict-IoU analysis, and same-hardware latency profiling. The experiments also separate solver network evaluations from cascade-head evaluations and show that solver acceleration and head distillation compress orthogonal sources of inference cost.

II. RELATED WORK

a) Chromosome detection and karyotyping: A metaphase image combines several difficulties rarely present together in natural-image detection: dozens of adjacent objects, homologous classes with similar morphology, scale variation between chromosome groups, staining shift, and incomplete boundaries caused by touching or overlap. Earlier work progressed from joint segmentation and classification [8], [9] to learned enumeration, straightening, and metaphase detection [10]–[12]. Recent 2026 studies confirm that the clinical workflow is moving toward direct detection from whole metaphase images: object detectors remove manual single-chromosome cropping [13], and integrated pipelines combine YOLO-based detection with pairing and anomaly screening [14]. However, those reports emphasize concordance, mAP50, or downstream workflow accuracy. These measures can remain high when a small chromosome has a several-pixel boundary error. Transfer-learning failure analysis likewise identifies overlap-related cropping artifacts as a persistent source of chromosome confusion [15]. KaryoFlow therefore targets the localization stage itself, with mAP over IoU 0.50:0.95, strict-IoU diagnostics, and measured inference cost.

Large-scale representation learning addresses a complementary limitation. CHROMA was pretrained on more than

four million chromosome images and improves robustness to rare abnormalities and challenging acquisition conditions [16]. Such single-chromosome representations reduce annotation burden after extraction, but do not resolve how densely overlapping candidates should move, retain identity, or be ranked in the preceding box-detection stage. This distinction motivates a detector specialized for metaphase box geometry rather than another isolated-chromosome classifier.

b) Generative detection and few-step integration: DiffusionDet [5] formulates object detection as iterative box denoising. FlowDet [17] and DeFloMat [18] replace DDPM-style dynamics with conditional flow matching or RF, while DiffuBox [19] studies a distinct 3D refinement setting. RF [6] and Flow Matching [20] provide our adopted continuous-time foundation, and DPM-Solver++ [7] supplies the multistep integration principle. Recent work has also explored flow-matching design and optimized priors for cellular-microscopy generation and medical augmentation [21], [22]; these image-generative settings are complementary to our discriminative transport of box sets. Existing studies emphasize the probability path or solver, but do not analyze the interaction between solver history and proposal renewal. KaryoFlow derives the data-prediction update for the RF box path and makes proposal identity an explicit condition of multistep inference. Diffusion in TMI has mainly addressed synthesis, segmentation, or robust classification [23]–[25], leaving few-step medical box-set prediction comparatively underexplored.

c) Localization quality in medical detection: Medical detectors must handle small targets, incomplete labels, acquisition shift, and structured false positives. Representative TMI studies address few-shot pulmonary lesions, heterogeneous lesion annotations, and relational transfer across domains [26]–[28], while context refinement, weak supervision, and imbalanced learning reduce task-specific false positives [29]–[31]. For general detection, IoU-Net predicts localization confidence for proposal selection [32], and GFL and VarifocalNet combine class and localization quality in dense heads [33], [34]. LQCR inherits the probability-ranking motivation but differs in causal scope: it is applied only after a generative detector has completed all trajectory updates, so solver history, coordinates, predicted classes, and proposal survival are fixed. This controlled boundary is particularly useful in chromosome images, where AP50 is nearly saturated but strict-IoU errors remain concentrated in small and overlapping instances. The emphasis on output reliability is also consistent with evidence that accurate medical networks can remain poorly calibrated [35].

III. METHOD

A. Box-Space Rectified Flow

a) Linear transport in chromosome-box space: Each chromosome is represented by a normalized center and size vector (c_x, c_y, w, h) . Training connects a Gaussian source box $\mathbf{x}_1$ to a matched chromosome box $\mathbf{x}_0$ through

$$\mathbf{x}_t = (1 - t) \mathbf{x}_0 + t \mathbf{x}_1, \quad t \in [0, 1], \quad (1)$$

whose target velocity is $\mathbf{x}_1 - \mathbf{x}_0$. Unlike image diffusion, only four box coordinates are transported; the microscopy features




Fig. 2. KaryoFlow from image to detections. (a) A metaphase image and its ResNet-50-FPN representation provide multi-scale image features. (b) A dense crop visualizes four-NFE box refinement from noise to localized chromosome candidates. (c) The detector emits the final quality-ranked chromosome detections. (d) One proposal is shown in normalized center-and-scale coordinates along the prescribed linear RF path. (e) At each evaluation, multi-scale RoIAlign feeds a cascade head with self-attention, DynamicConv, feed-forward refinement, AdaLN-Zero time conditioning, and the final-head quality branch. Endpoints and output boxes are empirical; intermediate boxes visualize the prescribed RF path rather than recorded hidden activations.

remain fixed conditioning information. This low-dimensional, straight path is a natural match for a metaphase spread, where approximately 46 boxes must move concurrently while preserving their individual identities.

b) Time conditioning: AdaLN-Zero [36] injects the time embedding into each proposal head through zero-initialized residual modulation. It is retained as a stable conditioning mechanism; controlled ablation does not support treating it as a separate accuracy contribution.

B. Stable Parameterization and Few-Step Integration

a) Few-step integration: Euler integration provides the first-order reference; it is inexpensive per step but uses only the current box estimate. Heun adds a predictor-corrector evaluation to reduce local integration error, so four reverse steps require seven network evaluations. DPM-Solver++ [7] instead reaches second order by reusing the previous target-box prediction, requiring only one new evaluation per step. We compare the solvers in this progression—Euler, Heun, then DPM-Solver++—and attribute solver gains only under matched checkpoints and step counts.

The model directly predicts $\mathbf{x}_0$ (the target box), and the update uses polynomial interpolation of the $\mathbf{x}_0$ history in t space (not the log-SNR λ space of VP-SDE). The RF-ODE in data-prediction form is $d\mathbf{x}/dt = (\mathbf{x} - \mathbf{x}_0(t))/t$; integrating exactly over $[t_n, t_{n+1}]$ with linear $\mathbf{x}_0(t)$ interpolation yields the update in (2)

$$\mathbf{x}_{t_{n+1}} = \frac{t_{n+1}}{t_n} \mathbf{x}_{t_n} + \left(1 - \frac{t_{n+1}}{t_n}\right) \mathbf{x}_0^{(n)} + \varphi_1 \mathbf{D}_1, \quad (2)$$

$$\varphi_1 = t_{n+1} \log \frac{t_n}{t_{n+1} - t_n + t_{n+1}}, \quad \mathbf{D}_1 = \frac{\mathbf{x}_0^{(n)} - \mathbf{x}_0^{(n-1)}}{t_n - t_{n-1}},$$

The first two terms move a proposal toward its current box estimate; the last term corrects that move using the change between the two most recent estimates. The initial noise sample remains fixed. Thus four steps require four rather than seven network evaluations. The endpoint is protected with $t_{n+1} > 10^{-7}$.

b) Why direct data prediction is numerically preferable here: For the path in (1), $\mathbf{x}_0 = \mathbf{x}_t - t\mathbf{v}$. Hence the two parameterizations carry the same information, but their pointwise squared errors satisfy

$$\mathcal{L}_v(t) = t^{-2} \mathcal{L}_{x_0}(t). \quad (3)$$

With the shifted schedule, velocity regression therefore amplifies gradient noise as $t \rightarrow 0$, where accurate terminal boxes matter most. Direct x_0 prediction removes this avoidable conditioning factor. We use this parameterization for its better-conditioned terminal objective, without claiming a test-set accuracy gain over algebraically equivalent velocity prediction.

c) Euler, Heun, and DPM-Solver++ comparison: Evaluating one RF checkpoint across the complete solver $\times$ step matrix (Table I) shows that solver choice is practically neutral at matched step counts. Four-step DPM-Solver++ reaches 0.8568 mAP, compared with 0.8566 for Heun and 0.8567 for Euler. Increasing DPM-Solver++ from one to four steps adds 0.0054 mAP. The solver is therefore an efficiency component rather than an accuracy component.

d) Conclusion: DPM-Solver++ is an efficiency component. On the same RF checkpoint at four steps, DPM-Solver++ and Heun differ by only +0.0002 aggregate mAP, while DPM-Solver++ uses four rather than seven network evaluations. The previously observed +0.0060 cross-checkpoint difference

TABLE I

LOCKED SOLVER $\times$ STEP DISENTANGLEMENT ON ONE RF CHECKPOINT.
VALUES ARE THREE-INFERENCE-SEED MEANS ON THE DATASET 2
HELD-OUT TEST; MATCHED-STEP SOLVER DIFFERENCES ARE NEGLIGIBLE.

Solver	Steps	NFE	mAP
Euler	1	1	0.8513
Euler	4	4	0.8567
Heun	4	7	0.8566
DPM-Solver++	1	1	0.8514
DPM-Solver++	4	4	0.8568

compared models trained under different configurations and cannot be attributed to the inference solver alone. We therefore claim matched accuracy at lower NFE, not a solver-induced accuracy improvement.

C. History-Consistent Proposal Integration

a) *Proposal identity is a premise of multistep history:*

Let $r_i^n = 1$ indicate that proposal row i is renewed after step n .

Proposition 1 (Identity consistency). *The multistep correction $\mathbf{D}_{1,i}$ in (2) is a consistent first-order approximation to a trajectory derivative only if the two states indexed by row i belong to the same proposal trajectory, or if the row-wise history is reset after replacement.*

Proof. Without renewal, differentiability of $\mathbf{x}_{0,i}(t)$ gives the standard secant limit

$$\mathbf{D}_{1,i} = \frac{\mathbf{x}_{0,i}^{(n)} - \mathbf{x}_{0,i}^{(n-1)}}{t_n - t_{n-1}} \rightarrow \frac{d\mathbf{x}_{0,i}}{dt}. \quad (4)$$

If $r_i^n = 1$, the numerator subtracts predictions from two independently initialized states. Its limit contains a finite cross-trajectory jump divided by $t_n - t_{n-1}$ and therefore is not the derivative along either trajectory. Resetting history removes the invalid correction. $\square$

We use the simpler alternative of disabling renewal after verifying sufficient proposal redundancy. This choice is well matched to metaphase images: the object count is tightly concentrated near 46, so $K = 500$ provides multiple candidates per chromosome, while maintaining identity avoids mixing adjacent, morphologically similar instances. In the locked Dataset 2 evaluation, disabling renewal at $K = 100$ costs 0.0169 mAP relative to retaining renewal, showing that history consistency cannot compensate for inadequate capacity.

D. Orthogonal Inference Compression

a) *Cascade-head distillation:* Solver NFE counts network calls, whereas each call evaluates H cascade heads. We distill a six-head teacher into a three-head student using head mapping $\{0 \rightarrow 0, 1 \rightarrow 2, 2 \rightarrow 5\}$ and feature/prediction supervision. At four solver steps this reduces cascade-head evaluations (CHE) from 24 to 12, while NFE remains four; the reduction is therefore 24-to-12 CHE, not NFE. The compressed model reaches 0.8595 versus 0.8598 and reduces RTX A6000 latency from 75.20 to 44.10 ms. Together with few-step integration,

this reduces latency from the 120.34-ms RF+Heun reference to 44.10 ms, a $2.73\times$ acceleration; the 0.0003 mAP cost is measured against the six-head teacher under the same fixed-checkpoint three-inference-seed test protocol.

E. Localization-Quality Calibrated Ranking

Chromosome classes can be predicted confidently even when their boxes are a few pixels too loose or shifted. This is especially harmful for small or touching chromosomes because a small displacement causes a large IoU change. For proposal feature F , correct-class event C , IoU U , and evaluation threshold τ , the true-positive probability separates as

$$P(C=1, U \geq \tau \mid F) = P(C=1 \mid F)P(U \geq \tau \mid C=1, F). \quad (5)$$

The second factor is absent from class-only ranking. LQCR estimates it with a two-layer quality branch $q = \sigma(g_\phi(F))$ attached only to the final cascade head. It adds 33,153 parameters. The backbone, FPN, proposal interaction, class head, and box head are frozen during branch fitting, so the experiment isolates output calibration from box generation.

For Hungarian-matched proposal i , let $m_i = 1$ and let u_i be the detached aligned IoU; unmatched proposals have $m_i = 0$. We train the scalar logit with the continuous target $y_i = m_i u_i$ and weighted binary cross-entropy

$$\mathcal{L}_q = \frac{\lambda_q}{\max(N_+, 1)} \sum_i w_i \ell_{\text{BCE}}(g_i, y_i), \quad (6)$$

$$w_i = m_i u_i + (1 - m_i) \alpha q_i^\gamma,$$

where $\lambda_q = 0.25$, $\alpha = 0.75$, and $\gamma = 2$. Detaching u_i prevents calibration from changing box geometry. Matched proposals learn their IoU; unmatched proposals learn zero. At output, class probability p is replaced by $s = pq^2$. A validation sweep over $\beta \in \{0, 0.25, 0.5, 1, 2\}$ selects the exponent 2; this learned-quality sweep is separate from the true-IoU oracle used only to quantify headroom.

Our strict final-only implementation keeps raw class scores in the solver, proposal renewal, Top- K pruning, and early stopping. Calibration is applied only when emitting detections, so boxes, predicted classes, and RF trajectories remain fixed at the candidate level; score-dependent NMS or max-detection selection may consequently choose a different final subset. An elementwise checkpoint audit confirms that all 590 tensors shared with the six-head detector are identical; the LQCR checkpoint contains only five additional quality-branch tensors. This design targets the observed chromosome regime in which AP50 is nearly saturated but small and overlapping instances lose recall at IoU 0.9. Figure 3 summarizes the principle, causal boundary, and fixed-seed controlled effect.

F. Top- K Proposal Pruning

After step 0 of inference, we prune proposals from 500 to K based on confidence scores; only the top- K proposals proceed through steps 1–3. Because pruning changes the proposal set, multistep history is reset before the retained proposals continue refinement.




Fig. 3. Effect and causal scope of localization-quality ranking. (a) On the fixed Dataset 2 RF-DPM++ checkpoint, quality-aware and class-only AP curves separate mainly at strict IoU thresholds. (b) Their threshold-wise difference reaches +2.78/+3.00 points at AP90/AP95, while AP50 changes by only −0.01 points. (c) The same candidate boxes and class scores are shown before class-only and LQCR ranking; localization quality promotes the tighter candidate and changes the retained top-2 subset without altering candidate geometry. Values in (a)–(b) are loaded from checkpoint-linked exports; (c) is a mechanism illustration.

TABLE II
DATASETS USED IN THIS PAPER. DATASET 1 IS CHROMOSOME20240904 (RST); DATASET 2 IS THE 24 CHROMOSOMES OBJECT DATASET.

Dataset	Train	Val	Test	Classes
Dataset 1	1,540	440	220	24
Dataset 2	3,500	500	1,000	24

IV. EXPERIMENTS

A. Experimental Setup

a) Datasets: Table II summarizes the two public chromosome datasets, referred to as Dataset 1 and Dataset 2. Dataset 1 is Chromosome20240904 (RST) [37], with 2,200 images; Dataset 2 is the 24 Chromosomes Object dataset [38], [39], with 5,000 images. Their training splits contain 1,540 and 3,500 images, respectively. Both have a median of 46 annotated instances per image, but their median relative box areas differ substantially (0.4% versus 0.8%). Under the COCO absolute-area definition, 23.0% of Dataset 1 instances are small, compared with only 0.4% in Dataset 2; correspondingly, 100% and 11.3% of test images contain at least one small instance. Moreover, 99.1% and 100% of test images contain at least one ground-truth box pair with IoU above 0.1. These statistics quantify the small, crowded regime and the scale shift used for cross-dataset replication. Both datasets use image-level random splitting. A single blood sample may yield multiple metaphase images, but patient identifiers are unavailable; thus patient-level separation cannot be verified. We analyze only the publicly released images and annotations and collect no additional human-subject data.

b) Architecture and training: Our base framework (KaryoFlow) uses ResNet-50 [40] + FPN [41] (256 channels, 4 levels), 500 proposals, 6 cascade transformer heads with deep supervision (5 auxiliary heads). Optimization: AdamW [42] (learning rate 5×10^{-5} , weight decay 10^{-4}), 5-epoch warmup + CosineAnnealing [43], 150 epochs. Loss: Focal [44] ($\lambda_{cls}=2$)

TABLE III
HELD-OUT-TEST PROPERTIES THAT MOTIVATE THE DETECTOR DESIGN. RELATIVE AREA IS BOX AREA DIVIDED BY IMAGE AREA; SMALL FOLLOWS THE COCO ABSOLUTE-AREA THRESHOLD ($< 32^2$ PIXELS); OVERLAP DENOTES AN IMAGE CONTAINING AT LEAST ONE PAIR OF GROUND-TRUTH BOXES WITH IoU > 0.1 .

Dataset	Median n	Med. rel. area	Small inst.	Images w/ small	Overlap
Dataset 1	46	0.4%	23.1%	100%	99.1%
Dataset 2	46	0.8%	0.4%	11.3%	100%

+ L1 ($\lambda=5$) + GIoU [45] ($\lambda=2$) with Hungarian matching [46]. Diffusion uses RF with shifted noise schedule (shift=3); default solver is DPM-Solver++ 4-step. For the isolated LQCR intervention, the trained six-head checkpoint initializes the model, all pre-existing parameters are frozen, and only the 33k-parameter quality branch is optimized for at most 12 epochs with AdamW learning rate 10^{-3} and cosine decay; the epoch-2 validation checkpoint is used. Thus this phase fits a ranking statistic but cannot alter the detector’s generated boxes or class logits.

c) Detection protocol: Boxes are represented in two spaces: image space (absolute pixel $xyxy$) for RoIAlign and NMS, and diffusion space ($cxcywh$ normalized to $[0, 1]$, mapped to $[-s, +s]$, $s=2$). The RF forward path is $x_t = (1-t)x_0 + t\varepsilon$ (DDPM baseline uses cosine schedule). At inference, 500 random-noise proposals are iteratively denoised; with time-ensemble, predictions from all steps are concatenated and deduplicated by per-class NMS (IoU 0.5). Evaluation uses COCO mAP@0.50:0.95 (10 IoU thresholds, maxDets=100). Validation data are used only for checkpoint and hyperparameter selection. Unless explicitly identified as a diagnostic oracle, every final detector comparison, locked hyperparameter confirmation, and accuracy coordinate in an efficiency plot is evaluated on the held-out test cohort (220 images for Dataset 1 and 1,000 for Dataset 2). Accuracy may be reproduced on different GPU models with fixed seeds; all latency and FPS measurements use the same RTX A6000

protocol on Ross.

d) Statistical considerations: Cross-seed experiments (3 seeds: 42, 123, 789) are available for Dataset 1. For Dataset 2, fixed-checkpoint inference seeds quantify stochastic proposal variation but do not replace independent training-seed uncertainty. The $+0.0060$ cross-checkpoint solver difference compares independently trained models and is not a solver-causal result; the same-checkpoint comparison is neutral. LQCR has clean three-seed replication on Dataset 1 and a fixed-checkpoint, final-only intervention on Dataset 2. The latter isolates ranking from training and trajectory changes; the former tests whether the effect reproduces across independent training runs.

e) Reproducibility: All reported comparisons retain their dataset split, random seed, checkpoint, and evaluation protocol. Controlled results are separated from exploratory runs and diagnostic oracles. For LQCR, we additionally verify that all 590 shared detector tensors are identical before fitting the quality branch and distinguish learned-quality sweeps from true-IoU oracle analyses. This protocol prevents cross-checkpoint or oracle quantities from entering model claims.

B. Main Results: RF vs DDPM

a) Dataset 2 ablation: Table IV reports a cumulative Dataset 2 ablation under the final evaluation protocol. The DDPM and RF+Heun rows establish the adopted generation foundation. One RF-DPM++ checkpoint then isolates the four-NFE DPM-Solver++ operating point and the final-only LQCR intervention. AdaLN-Zero is used throughout as time conditioning but is not counted as an independent contribution.

RF supplies the dominant historical improvement over the DDPM baseline. DPM-Solver++ retains accuracy while reducing Heun’s seven evaluations to four; same-checkpoint solver tests do not support an additional precision claim. LQCR then improves mAP by 0.0070 in the fixed-checkpoint three-inference-seed test without changing the checkpoint, candidate coordinates, class logits, renewal decisions, or trajectory. Its largest gains occur at strict IoU thresholds (AP90/AP95 $+0.0278/+0.0300$), consistent with the localization-aware ranking mechanism rather than a coordinate change.

Head distillation is reported separately in the deployment benchmark because the three-head student was trained only on Dataset 2. Validation-only and single-seed exploratory directions are excluded from this table.

b) Dataset 1: RF vs DDPM: On the Dataset 1 held-out test split (3 seeds), RF+Heun reaches 0.7367 ± 0.0015 mAP versus 0.7187 ± 0.0031 for DDPM, an absolute gain of 0.0180. On Dataset 2, the unified held-out test comparison gives 0.8566 for RF+Heun and 0.8033 for DDPM, an absolute gain of 0.0533. These comparisons use matched dataset splits and are therefore preferred over the older $+0.0820$ validation estimate, which mixed different training schedules and is excluded from current claims.

C. SOTA Comparison

Table VI uses the locked final protocol for both chromosome datasets. The cohorts differ strongly in relative object

scale, yet the full KaryoFlow model with LQCR obtains the highest mAP among the evaluated methods on each: 0.7480 on Dataset 1 and 0.8668 on Dataset 2. The absolute margins over the strongest listed alternatives are 0.0161 and 0.0019 (relative 2.2% and 0.2%), respectively. The result is not caused by a larger backbone: KaryoFlow uses ResNet-50-FPN, whereas RTMDet-L uses CSPNeXt-L and DINO uses multi-scale deformable attention [3], [47]. We therefore describe these as best results within the checkpoint-linked comparison, rather than claiming an exhaustive cross-study SOTA.

a) Small-object performance: Dataset 1 is the small-object-heavy cohort. The three-seed LQCR mean is 0.7480 ± 0.0017 mAP and 0.5147 ± 0.0090 AP_S, improving the uncalibrated KaryoFlow means by 0.0067 and 0.0057 and exceeding the strongest listed AP_S comparator by 0.0142. On Dataset 2, LQCR leads overall mAP but remains below DINO on AP_S. This pattern supports the intended role of quality calibration—better aggregate ordering across strict IoU thresholds—without claiming universal superiority at every scale.

b) Class- and scale-resolved analysis: The held-out tests show that AP50 is already near saturation, whereas the quality-ranking gain concentrates at AP90/AP95. The remaining headroom therefore lies in precise coordinate refinement and localization-aware ordering rather than simply producing more low-quality proposals. This pattern is especially relevant to the smaller object distribution in Dataset 1.

c) Causal interpretation of controlled comparisons: The same-checkpoint solver comparison is accuracy-neutral. The retained positive intervention is LQCR: enabling its score only at the final output changes the three-inference-seed mean mAP from 0.8598 to 0.8668 on exactly the same checkpoint and detections before ranking. This paired construction isolates decision-stage calibration from generation, interaction, and numerical integration effects. We report this as a fixed-checkpoint causal intervention, not as a cross-seed training gain; Dataset 1 additionally supplies independent training-seed replication.

The monotone mAP trend confirms the validation-selected $\beta = 2$. AP90 and AP95 rise at every point in the locked sweep, matching the threshold-factorization prediction.

1) Scale-Resolved Cross-Dataset Comparison: Figure 4 replaces selected visual cases with complete-set evidence. Dataset 1 has a substantially smaller relative-box distribution than Dataset 2 (median 0.4% versus 0.8%). On the Dataset 1 test split, KaryoFlow+LQCR reaches AP_S = 0.5147 ± 0.0090 over three seeds, a 0.0057 gain over the uncalibrated model and 0.0142 above the strongest listed comparator. On Dataset 2, LQCR reaches AP_S = 0.5663 but remains below DINO R50 (0.5953). The result supports scale-dependent small-object competitiveness rather than a universal advantage. Because COCO’s small-area threshold is absolute, raw AP_S levels are interpreted only within each dataset; the relative-area CDF is used to characterize the cross-dataset scale shift.

D. Solver Analysis

a) DPM-Solver++ step ablation: DPM-Solver++ improves from 0.8514 at one step to 0.8568 at four steps, consistent with rapid convergence along the short RF box trajectories.

TABLE IV

EVIDENCE-CONTROLLED DATASET 2 ABLATION. THE FINAL TWO ROWS USE ONE FIXED RF-DPM++ LINEAGE AND REPORT THREE INFERENCE-SEED TEST MEANS; LQCR CHANGES ONLY OUTPUT RANKING.

Experiment	Solver	Steps	NFE	mAP	AP50	AP75	AP _S	AP _M	AP _L
DDPM baseline (Euler 1-step)	Euler	1	1	0.8033	0.9714	0.9381	0.5323	0.8006	0.8221
RF+Heun (KaryoFlow)	Heun	4	7	0.8566	0.9878	0.9691	0.5827	0.8541	0.8496
DPM-Solver++	DPM++	4	4	0.8598	0.9883	0.9709	0.5676	0.8572	0.8876
+LQCR (final-only)	DPM++	4	4	0.8668	0.9882	0.9714	0.5663	0.8640	0.8890

TABLE V

CORE INTERVENTIONS ON BOTH HELD-OUT TEST COHORTS. DATASET 1 LQCR USES THREE INDEPENDENTLY TRAINED CHECKPOINTS; FIXED-CHECKPOINT COMPARISONS REPORT THE MEAN OVER INFERENCE SEEDS 42, 123, AND 789.

Intervention	D1 Δ mAP	D2 Δ mAP	Deployment effect	Role and boundary
LQCR	+0.0067	+0.0070	+3.3% latency	Final-only ranking; candidate boxes and trajectories are unchanged
Renewal off ($K = 500$)	-0.0016	-0.0004	-2.4% latency	Preserves DPM++ history identity; not claimed as an accuracy gain

TABLE VI

LOCKED COMPARISON ON TWO CHROMOSOME DATASETS. DATASET 1 KARYOFlow+LQCR IS THE CLEAN THREE-TRAINING-SEED MEAN; DATASET 2 AND THE OTHER ROWS ARE FIXED-CHECKPOINT THREE-INFERENCE-SEED TEST RE-EVALUATIONS. BOLD DENOTES THE BEST mAP WITHIN THIS COMPARISON, NOT AN EXHAUSTIVE CROSS-PAPER SOTA CLAIM.

Method	Backbone	Dataset 1				Dataset 2			
		mAP	AP50	AP75	AP _S	mAP	AP50	AP75	AP _S
KaryoFlow+LQCR	ResNet-50	0.7480	0.9323	0.8330	0.5147	0.8668	0.9882	0.9714	0.5663
KaryoFlow	ResNet-50	0.7402	0.9313	0.8244	0.5026	0.8598	0.9883	0.9709	0.5676
DINO R50	ResNet-50	0.7248	0.9341	0.8135	0.4853	0.8649	0.9890	0.9728	0.5953
RTMDet-L	CSPNeXt-L	0.7319	0.9344	0.8387	0.4803	0.8609	0.9890	0.9706	0.5725
Cascade R-CNN	ResNet-50	0.7235	0.9221	0.8307	0.5005	0.8527	0.9850	0.9670	0.5724
DiffusionDet	ResNet-50	0.7136	0.9078	0.7975	0.4618	0.8033	0.9714	0.9381	0.5323
YOLOX-S	CSPDarkNet-S	0.5811	0.9177	0.6920	0.3298	0.7951	0.9854	0.9427	0.4861

TABLE VII

LOCKED LEARNED-QUALITY EXPONENT CONFIRMATION ON DATASET 2. VALUES ARE THREE-INFERENCE-SEED MEANS ON THE HELD-OUT TEST. VALIDATION SELECTED $\beta = 2$; $\beta = 0$ DISABLES QUALITY RANKING. THESE ARE LEARNED- q RESULTS, NOT GROUND-TRUTH-IOU ORACLE RANKINGS.

β	mAP	Δ mAP	AP90	AP95
0	0.8598	—	0.6791	0.2053
0.25	0.8622	+0.0024	0.6890	0.2153
0.5	0.8638	+0.0040	0.6946	0.2221
1	0.8654	+0.0056	0.7014	0.2290
2	0.8668	+0.0070	0.7069	0.2353

b) *DPM-Solver++ vs Heun at matched NFE*: At four matched steps on the same checkpoint, DPM-Solver++ gives 0.8568, Heun 0.8566, and Euler 0.8567. The largest solver gap is only 0.0002, well below the observed inference-seed variation. We therefore attribute no accuracy gain to the off-the-shelf solver; its value is reducing the four-step Heun computation from seven to four network evaluations.

E. FPS / Latency Benchmark

Table VIII and Figure 1 report the speed-accuracy trade-off at 512×512 input resolution. Latency is CUDA-event model-forward time and excludes image decoding, host pre-


TABLE VIII

SAME-HARDWARE MODEL-FORWARD LATENCY BENCHMARK ON DATASET 2 AT 512×512 RESOLUTION (RTX A6000, BATCH SIZE 1, 10 WARM-UP AND 500 TIMED CUDA-EVENT ITERATIONS). ACCURACY FOLLOWS THE LOCKED FINAL PROTOCOL; ONLY LATENCY IS HARDWARE-DEPENDENT.

Model	Solver	NFE	Latency (ms)	FPS	mAP
RF+Heun	Heun	7	120.34	8.3	0.8566
DPM-Solver++	DPM++	4	75.20	13.3	0.8598
KaryoFlow+LQCR	DPM++	4	77.68	12.9	0.8668
DPM-Solver++ +Top-K ($K=200$)	DPM++	4	70.57	14.2	0.8590
DPM-Solver++ +Top- K ($K=100$)	DPM++	4	68.90	14.5	0.8480
DINO R50	—	1	34.33	29.1	0.8649
RTMDet-L	—	1	33.83	29.6	0.8609
Cascade R-CNN	—	1	20.76	48.2	0.8527
DiffusionDet	Euler	1	24.09	41.5	0.8033
YOLOX-S	—	1	10.04	99.7	0.7951

processing, and visualization. The unpruned $K = 500$ point and the $K = 200/100$ pruning points share one checkpoint and evaluation protocol. DPM-Solver++ with Top- K pruning reaches a latency compatible with interactive use, while single-shot detectors (Cascade R-CNN, YOLOX-S) trade accuracy for speed.

DPM-Solver++ +Top- K ($K=200$) reaches 70.57 ms / 14.2 FPS at mAP 0.8590. The complete KaryoFlow+LQCR model reaches 77.68 ms / 12.9 FPS at mAP 0.8668; its 2.48 ms



(b)

Measure	Dataset 1	Dataset 2
Small-instance share	23.0%	0.4%
KaryoFlow AP_5	0.5090	0.5728
KaryoFlow+LQCR AP_5	0.5147 $\pm$ 0.0090	0.5661
LQCR ΔAP_5	+0.0057	-0.0068
Best comparator	Cascade R-CNN	DINO R50
Gap to comparator in AP_5	+0.0142	-0.0288
Overall mAP rank	1st	1st

AP_5 is interpreted within each cohort; Dataset 1 LQCR is the 3-seed mean.

Fig. 4. Object-scale shift and small-object detection across two chromosome cohorts. (a) The empirical CDF of relative ground-truth box area over all test annotations shows that Dataset 1 chromosomes are markedly smaller than those in Dataset 2. (b) A compact scale-response table contrasts the small-instance prevalence, LQCR effect, strongest comparator, and overall mAP rank. AP_5 is interpreted within, not across, datasets; Dataset 1 LQCR is the clean three-seed mean.

(3.3%) overhead over uncalibrated DPM-Solver++ buys the highest accuracy in the comparison. The cascade head dominates 90%+ of latency; the backbone+neck is a minor cost (~ 5.6 ms). These speed-linked re-evaluations differ by at most 0.0030 from the canonical training-time point estimates in Tables IV and VI, which is consistent with random proposal initialization. The RF+Heun row uses the same Dataset 2 A1 checkpoint for both axes.

The efficiency analysis separates three counts that are often conflated: DPM-Solver++ reduces solver NFE ($7 \rightarrow 4$), and head distillation reduces CHE ($24 \rightarrow 12$ at four NFE). Renewal removal does not change either count; it makes multistep history coherent and eliminates proposal replacement overhead.

V. DISCUSSION

KaryoFlow’s components address different consequences of metaphase-image structure. The low-dimensional linear path and direct box prediction reduce optimization burden when annotated cohorts are modest. Proposal identity matters because dozens of adjacent candidates evolve simultaneously, while the nearly fixed chromosome count permits a redundant proposal pool that makes renewal unnecessary at the retained operating point. LQCR then addresses the remaining decision bottleneck: class confidence is already sufficient at IoU 0.5, but it does not identify which small or overlapping boxes will survive stricter localization thresholds. The concentration of LQCR gains at AP_{90} and AP_{95} is consistent with this task-derived mechanism.

The efficiency results reveal complementary deployment controls rather than a single opaque speedup. DPM-Solver++ reduces solver evaluations, distillation reduces cascade-head evaluations, and Top- K reduces proposals after the first step. This decomposition enables an accuracy-oriented setting (0.8668 mAP, 77.68 ms) or a compressed setting (0.8595 mAP, 44.10 ms) without changing the central trajectory-ranking formulation.

The study has three boundaries. First, both cohorts are public chromosome datasets and lack patient identifiers, so

patient-level leakage and clinical workflow outcomes cannot be assessed. Second, LQCR is replicated across three clean Dataset 1 training seeds and confirmed with a locked final-only intervention on Dataset 2, but the latter still uses one trained checkpoint. Natural-image validation is required to establish broader generality. Third, history-consistent inference relies on proposal redundancy: the observed $K = 100$ failure should be treated as an operating constraint, not hidden by the high- K result. These limits do not affect the box-path derivations, but they define the present empirical scope.

VI. CONCLUSION

KaryoFlow aligns trajectory learning, multistep state identity, and final localization-aware ranking for dense chromosome detection. RF supplies the few-step generative foundation; direct data prediction provides a well-conditioned terminal objective; and the proposal-identity analysis makes DPM++ history mathematically coherent at the retained proposal count. The independent decision contribution, LQCR, improves test mAP by 0.0067 across three Dataset 1 seeds and by 0.0070 in the Dataset 2 fixed-checkpoint intervention. On Dataset 2, AP_{90}/AP_{95} improve by 0.0278/0.0300 without changing candidate boxes before ranking. Together with orthogonal inference compression, these results establish a competitive and interpretable generative detector for chromosome karyotyping.

REFERENCES

- [1] Z. Cai and N. Vasconcelos, “Cascade R-CNN: Delving into high quality object detection,” in *Proc. IEEE Conf. Comput. Vis. Pattern Recognit. (CVPR)*, 2018, pp. 6154–6162.
- [2] Z. Ge, S. Liu, F. Wang, Z. Li, and J. Sun, “YOLOX: Exceeding YOLO series in 2021,” arXiv:2107.08430, 2021. [Online]. Available: <https://arxiv.org/abs/2107.08430>
- [3] H. Zhang, F. Li, S. Liu, L. Zhang, H. Su, J. Zhu, L. M. Ni, and H.-Y. Shum, “DINO: DETR with improved denoising anchor boxes for end-to-end object detection,” in *Int. Conf. Learn. Represent. (ICLR)*, 2023.
- [4] X. Zhu, W. Su, L. Lu, B. Li, X. Wang, and J. Dai, “Deformable DETR: Deformable transformers for end-to-end object detection,” in *Int. Conf. Learn. Represent. (ICLR)*, 2021.

- [5] S. Chen, P. Sun, Y. Song, and P. Luo, "DiffusionDet: Diffusion model for object detection," in *Proc. IEEE/CVF Int. Conf. Comput. Vis. (ICCV)*, 2023, pp. 1980–1989.
- [6] X. Liu, C. Gong, and Q. Liu, "Flow straight and fast: Learning to generate and transfer data with rectified flow," in *Int. Conf. Learn. Represent. (ICLR)*, 2023.
- [7] C. Lu, Y. Zhou, F. Bao, J. Chen, C. Li, and J. Zhu, "DPM-Solver++: Fast solver for guided sampling of diffusion probabilistic models," *Machine Intelligence Research*, 2025.
- [8] W. C. Schwartzkopf, A. C. Bovik, and B. L. Evans, "Maximum-likelihood techniques for joint segmentation–classification of multispectral chromosome images," *IEEE Transactions on Medical Imaging*, vol. 24, no. 12, pp. 1593–1610, 2005.
- [9] Y. Qin, J. Wen, H. Zheng, X. Huang, J. Yang, N. Song, Y.-M. Zhu, L. Wu, and G.-Z. Yang, "Varifocal-Net: A chromosome classification approach using deep convolutional networks," *IEEE Transactions on Medical Imaging*, vol. 38, no. 11, pp. 2569–2581, 2019.
- [10] L. Xiao, C. Luo, T. Yu, Y. Luo, M. Wang, F. Yu, Y. Li, C. Tian, and J. Qiao, "DeepACEv2: Automated chromosome enumeration in metaphase cell images using deep convolutional neural networks," *IEEE Transactions on Medical Imaging*, vol. 39, no. 12, pp. 3920–3932, 2020.
- [11] J. Li, S. Zheng, Z. Shui, S. Zhang, L. Yang, Y. Sun, Y. Zhang, H. Li, Y. Ye, P. M. A. van Ooijen, K. Li, and L. Yang, "Masked conditional variational autoencoders for chromosome straightening," *IEEE Transactions on Medical Imaging*, vol. 43, no. 1, pp. 216–228, 2024.
- [12] C.-E. Kuo, J.-Z. Li, J.-J. Tseng, F.-C. Lo, M.-J. Chen, and C.-H. Lu, "ChromosomeNet: Deep learning-based automated chromosome detection in metaphase cell images," *IEEE Open Journal of Engineering in Medicine and Biology*, vol. 6, pp. 227–236, 2024.
- [13] S. Matsumoto, M. Sojo, K. Oshima, K. Nishikawa, K. Okubo *et al.*, "Detection and classification of chromosomes with sister chromatid cohesion defects using object detection models," *Scientific Reports*, vol. 16, p. 13719, 2026.
- [14] C. Rosas-Alatríste, N. O. Rodríguez-Rodríguez, A. J. Argüelles-Cruz, A. Rosas-Cabral, C. Yáñez-Márquez *et al.*, "Automated karyotyping and structural anomaly detection through a hybrid multi-stage deep learning framework integrating chromosome detection, pairwise classification, and autoencoder-based analysis," *Scientific Reports*, vol. 16, p. 22469, 2026.
- [15] T. Chen, C. Xie, W. Zhang *et al.*, "ImageNet pre-training and two-step transfer learning in chromosome image classification," *Scientific Reports*, vol. 16, p. 7572, 2026.
- [16] C. Yang *et al.*, "A comprehensive foundation model for generalizable cytogenetics in precision oncology with CHROMA," *npj Precision Oncology*, vol. 10, p. 187, 2026.
- [17] E. Baty, C. P. Bridges, and S. Hadfield, "FlowDet: Unifying object detection and generative transport flows," arXiv:2512.16771, 2025. [Online]. Available: <https://arxiv.org/abs/2512.16771>
- [18] H. Lee, C. Lee, N. Seo, J. S. Lim, and H. Hong, "DeFloMat: Detection with flow matching for stable and efficient generative object localization," arXiv:2512.22406, 2025. [Online]. Available: <https://arxiv.org/abs/2512.22406>
- [19] X. Chen, Z. Liu, K. Z. Luo, S. Datta, A. Polavaram, Y. Wang *et al.*, "DiffuBox: Refining 3D object detection with point diffusion," in *Adv. Neural Inf. Process. Syst. (NeurIPS)*, 2024.
- [20] Y. Lipman, R. T. Q. Chen, H. Ben-Hamu, M. Nickel, and M. Le, "Flow matching for generative modeling," in *Int. Conf. Learn. Represent. (ICLR)*, 2023.
- [21] C. Jones, E. Noutahi, J. Hartford, and C. Eastwood, "Elucidating the design space of flow matching for cellular microscopy," arXiv:2603.26790, 2026. [Online]. Available: <https://arxiv.org/abs/2603.26790>
- [22] F. Nützel, M. Dombrowski, and B. Kainz, "Flow matching with optimized subclass priors for medical image augmentation," arXiv:2605.16469, 2026. [Online]. Available: <https://arxiv.org/abs/2605.16469>
- [23] M. Tivnan, J. Teneggi, T.-C. Lee, R. Zhang, K. Boedeker, L. Cai, G. J. Gang, J. Sulam, and J. W. Stayman, "Fourier diffusion models: A method to control MTF and NPS in score-based stochastic image generation," *IEEE Transactions on Medical Imaging*, vol. 44, no. 9, pp. 3694–3704, 2025.
- [24] H. Yang, W. Hua, Z. Xu, and J. Sun, "Domain-generalized discrete diffusion model for cross-domain medical image segmentation," *IEEE Transactions on Medical Imaging*, vol. 44, no. 11, pp. 4131–4142, 2025.
- [25] X. Shen, H. Huang, B. Nichyporuk, and T. Arbel, "Improving robustness and reliability in medical image classification with latent-guided diffusion and nested-ensembles," *IEEE Transactions on Medical Imaging*, vol. 44, no. 12, pp. 4890–4902, 2025.
- [26] W. Yang, S. Gao, H. Zhang, H. Yu, M. Xu, P. Chong, W. Zhang, H. Wang, W. Zhang, and A. Qian, "PtbNet: Based on local few-shot classes and small objects to accurately detect PTB," *IEEE Transactions on Medical Imaging*, vol. 43, no. 12, pp. 4331–4343, 2024.
- [27] K. Yan, J. Cai, Y. Zheng, A. P. Harrison, D. Jin, Y. Tang, Y. Tang, L. Huang, J. Xiao, and L. Lu, "Learning from multiple datasets with heterogeneous and partial labels for universal lesion detection in CT," *IEEE Transactions on Medical Imaging*, vol. 40, no. 10, pp. 2759–2770, 2021.
- [28] C. Chen, J. Wang, J. Pan, C. Bian, and Z. Zhang, "GraphSKT: Graph-guided structured knowledge transfer for domain adaptive lesion detection," *IEEE Transactions on Medical Imaging*, vol. 42, no. 2, pp. 507–518, 2023.
- [29] Y. Ren, X. Liu, J. Ge, Z. Liang, X. Xu, L. J. Grimm, J. Go, J. R. Marks, and J. Y. Lo, "Ipsilateral lesion detection refinement for tomosynthesis," *IEEE Transactions on Medical Imaging*, vol. 42, no. 10, pp. 3080–3090, 2023.
- [30] R. Wang, B. Chen, D. Meng, and L. Wang, "Weakly supervised lesion detection from fundus images," *IEEE Transactions on Medical Imaging*, vol. 38, no. 6, pp. 1501–1512, 2019.
- [31] S.-H. Bae and K.-J. Yoon, "Polyp detection via imbalanced learning and discriminative feature learning," *IEEE Transactions on Medical Imaging*, vol. 34, no. 11, pp. 2379–2393, 2015.
- [32] B. Jiang, R. Luo, J. Mao, T. Xiao, and Y. Jiang, "Acquisition of localization confidence for accurate object detection," in *Proc. Eur. Conf. Comput. Vis. (ECCV)*, 2018, pp. 784–799.
- [33] X. Li, W. Wang, L. Wu, S. Chen, X. Hu, J. Li, J. Tang, and J. Yang, "Generalized focal loss: Learning qualified and distributed bounding boxes for dense object detection," in *Adv. Neural Inf. Process. Syst. (NeurIPS)*, vol. 33, 2020, pp. 21 002–21 012.
- [34] H. Zhang, Y. Wang, F. Dayoub, and N. Sunderhauf, "VarifocalNet: An IoU-aware dense object detector," in *Proc. IEEE/CVF Conf. Comput. Vis. Pattern Recognit. (CVPR)*, 2021, pp. 8514–8523.
- [35] A. Mehrtash, W. M. Wells, C. M. Tempny, P. Abolmaesumi, and T. Kapur, "Confidence calibration and predictive uncertainty estimation for deep medical image segmentation," *IEEE Transactions on Medical Imaging*, vol. 39, no. 12, pp. 3868–3878, 2020.
- [36] P. Dhariwal and A. Nichol, "Diffusion models beat GANs on image synthesis," in *Adv. Neural Inf. Process. Syst. (NeurIPS)*, 2021.
- [37] South China Normal University, "RST (chromosome20240904) dataset," Roboflow Universe, 2024, accessed Jul. 2026. [Online]. Available: <https://universe.roboflow.com/south-china-normal-university-imqzk/rst>
- [38] J.-J. Tseng, C.-H. Lu, J.-Z. Li, H.-Y. Lai, M.-H. Chen, F.-Y. Cheng, and C.-E. Kuo, "An open dataset of annotated metaphase cell images for chromosome identification," *Scientific Data*, vol. 10, p. 104, 2023.
- [39] C.-H. Lu, C.-E. Kuo, and J.-J. Tseng, "CIL:54816 - 24 chromosomes object dataset," Cell Image Library, 2022. [Online]. Available: <https://doi.org/10.7295/W9CIL54816>
- [40] K. He, X. Zhang, S. Ren, and J. Sun, "Deep residual learning for image recognition," in *Proc. IEEE Conf. Comput. Vis. Pattern Recognit. (CVPR)*, 2016, pp. 770–778.
- [41] T.-Y. Lin, P. Dollár, R. Girshick, K. He, B. Hariharan, and S. Belongie, "Feature pyramid networks for object detection," in *Proc. IEEE Conf. Comput. Vis. Pattern Recognit. (CVPR)*, 2017, pp. 936–944.
- [42] I. Loshchilov and F. Hutter, "Decoupled weight decay regularization," in *Int. Conf. Learn. Represent. (ICLR)*, 2019.
- [43] —, "SGDR: Stochastic gradient descent with warm restarts," in *Int. Conf. Learn. Represent. (ICLR)*, 2017.
- [44] T.-Y. Lin, P. Goyal, R. Girshick, K. He, and P. Dollár, "Focal loss for dense object detection," in *Proc. IEEE/CVF Int. Conf. Comput. Vis. (ICCV)*, 2017, pp. 2999–3007.
- [45] H. Rezatofighi, N. Tsoi, J. Gwak, A. Sadeghian, I. Reid, and S. Savarese, "Generalized intersection over union: A metric and a loss for bounding box regression," in *Proc. IEEE Conf. Comput. Vis. Pattern Recognit. (CVPR)*, 2019, pp. 658–666.
- [46] H. W. Kuhn, "The hungarian method for the assignment problem," *Naval Research Logistics Quarterly*, vol. 2, no. 1-2, pp. 83–97, 1955.
- [47] C. Lyu, W. Zhang, H. Huang, Y. Zhou, Y. Wang, Y. Liu, S. Zhang, and K. Chen, "RTMDet: An empirical study of designing real-time object detectors," arXiv:2212.07784, 2022.